
COMP3109

Assignment 3

Compiler Optimizer (10 marks)

This third assignment is a **group assignment**. You can work in groups of up to three. You can work in the same groups as for assignment 2, or change groups. At least one group member must demonstrate the group's assignment at the tutorial in Week 13. The source code must be submitted via eLearning on or before **Friday October 31**. Your assignment will not be assessed unless all of the following criteria are met:

1. Submit a tarball of your source code to eLearning. Documentation should be in form of comments. Please provide plenty of them.
2. Hand in a signed academic honesty form in the tutorial before, or scan it and submit it in the tarball.
3. You must include the `optimize.sh` and `README` files as described in this assignment spec.
4. Your `README` file must list 1) the names of all group members, 2) the names and version numbers of all languages, compilers, and tools used.
5. Your assignment must build cleanly with the listed compilers and tools. No errors, no warnings.
6. Your tarball must *not* contain any platform specific binary build products, such as executables or object files. Cross platform products like `.jar` files and portable code generated from other tools is ok.

In this assignment you will implement an optimizer for the intermediate language from the previous assignment. Your optimizer should read intermediate code from a file, construct a control-flow graph (CFG) for each function, apply optimizations on the individual CFGs, and finally write the optimized intermediate code to another file.

The CFG of your optimizer should consist of control-flow edges and basic blocks, which contain intermediate code instructions. Be careful to take the other tasks of this assignment into consideration when designing the corresponding data structures. A flexible design will be of help during the implementation of the subsequent optimizations.

Once the CFG is available, several optimizations should be performed. The first optimization eliminates unreachable code. The optimization should be implemented in two phases: first perform a depth-first traversal of the CFG, building a set of all reachable blocks. Remove the unreachable blocks in a second phase.

The second optimization eliminates dead code, meaning instructions that produce values that are not used. Dead code elimination uses data-flow analysis to determine which results are used by a program and which can be eliminated without affecting the result of the program. The optimization consists of two phases. First, a data-flow analysis determines which registers hold used values. You should

implement this with an iterative solver that computes a fixed-point solution of the data-flow analysis problem. Remove the dead instructions in a second phase.

The final optimization eliminates redundant load instructions (ld). You should define a data-flow analysis that determines whether the value read by each load instruction is already available in some existing register. If a value is already available then it can be copied into the destination register, rather than re-loading it from memory. Unused load instructions can then be removed by the dead code eliminator you defined previously.

Your optimizer should be able to apply a single optimization pass to the input program, as well as apply all passes repeatedly until no further improvements can be achieved. The optimizer should work with any syntactically correct intermediate code. You should not limit your tests to files generated by the front-end developed for Assignment 2.

Intermediate Code

The grammar of the intermediate language is defined below (it is unchanged with regard to Assignment 2).

Program Structure

A program is a list of functions, where a function is a list consisting of the function's name followed by a lists of basic blocks. Each basic block in turn is a list consisting of a block number and a list of instructions.

Instructions

The intermediate code supports 13 different instructions:

- Load constant (`lc`)
Copy a constant into a register.
Example:
`(lc r5 5)` ... copy the value 5 into register 5
- Load instructions (`ld`)
Read the value of a variable and copy its value into a register.
Example:
`(ld r5 a)` ... read the value of a and copy it into register 5
- Store instructions (`st`)
Write the value of a register to a variable.
Example:
`(st b r5)` ... read the value of register 5 and assign it to variable b
- Arithmetic instructions (`add`, `sub`, `mul`, `div`)
Perform the respective arithmetic operations on registers.
Example:
`(add r3 r4 r5)` ... read register 4 and 5 and write the sum into register 3
- Comparison instructions (`lt`, `gt`, `feq`)
Perform a comparison and write the value 0 (false) or 1 (true) into a register.
Example:
`(eq r3 r4 r5)` ... compare register 4 and 5 and write the result into register 3
- Branch instructions (`br`)
Perform a conditional branch depending on a register value. If the value of the register is non-zero the execution continues at the basic block whose number is specified by the second operand. The execution, otherwise, resumes at the basic block specified by the third operand.
Example:
`(br r3 1 2)` ... depending on register 3 branch to basic block 1 or 2
- Return instructions (`ret`)
Return the value of a register from a function.
Example:
`(ret r3)` ... exit the current function and return the value of register 3

- Call instructions (`call`)

Perform a function call. The second argument specifies the function to be invoked, followed by a list of registers whose values should be passed as function arguments. The return value of the called function is written into the register specified as the first operand.

Example:

(`call r1 factorial r3`) ... invoke function `factorial` and pass the value of register 3 as its first (and only) argument; write the return value into register 1

Intermediate Code Grammar

```

<program>      ::= ( <functions> )

<functions>    ::= ε
<functions>    ::= <function> <functions>

<function>     ::= ( <ID> <arguments> <blocks> )

<arguments>    ::= ( <id_list> )

<id_list>      ::= ε
<id_list>      ::= <ID> <id_list>

<blocks>       ::= ( <NUM> instructions )
<blocks>       ::= ( <NUM> instructions ) blocks

<instructions> ::= <instruction>
<instructions> ::= <instruction> <instructions>

<instruction>  ::= ( lc <REG> <NUM> )
<instruction>  ::= ( ld <REG> <ID> )
<instruction>  ::= ( st <ID> <REG> )
<instruction>  ::= ( add <REG> <REG> <REG> )
<instruction>  ::= ( sub <REG> <REG> <REG> )
<instruction>  ::= ( mul <REG> <REG> <REG> )
<instruction>  ::= ( div <REG> <REG> <REG> )
<instruction>  ::= ( lt <REG> <REG> <REG> )
<instruction>  ::= ( gt <REG> <REG> <REG> )
<instruction>  ::= ( eq <REG> <REG> <REG> )
<instruction>  ::= ( br <REG> <NUM> <NUM> )
<instruction>  ::= ( ret <REG> )
<instruction>  ::= ( call <REG> <ID> <reg_list> )

<reg_list>     ::= ε
<reg_list>     ::= <REG> <reg_list>

<NUM>          ::= -?[0-9]+
<ID>           ::= [a-zA-Z][a-zA-Z0-9]*
<REG>          ::= r[1-9][0-9]*

```

Example

The unoptimized intermediate code of an example factorial program might initially look like this:

```
( (factorial (n)
  (0  (ld r1 n)
      (lc r2 0)
      (eq r3 r1 r2)
      (st cond r3)
      (ld r4 cond)
      (br r4 1 2) )
  (1  (lc r5 1)
      (st tmp r5)
      (ld r6 tmp)
      (ret r6) )
  (2  (ld r7 n)
      (lc r8 1)
      (sub r9 r7 r8)
      (st tmp r9)
      (ld r10 tmp)
      (call r11 factorial r10)
      (ld r12 n)
      (mul r13 r11 r12)
      (st tmp r13)
      (ld r14 tmp)
      (ret r14) ) )
(main (n)
  (0  (ld r1 n)
      (call r2 factorial r1)
      (st tmp r2)
      (ld r3 tmp)
      (ret r3) ) ) )
```

Your optimizer should produce an optimized program similar to the one below. Your optimizer does not need to produce *exactly* the same program. Minor differences in register numbering, instruction order, and instruction count are ok.

```
( (factorial (n)
  (0  (ld r1 n)
      (lc r2 0)
      (eq r3 r1 r2)
      (st cond r3)
      (ld r4 cond)
      (br r4 1 2) )
  (1  (lc r5 1)
      (st tmp r5)
      (ld r6 tmp)
      (ret r6) )
```

```
(2  (lc r8 1)
    (sub r9 r1 r8)
    (st tmp r9)
    (ld r10 tmp)
    (call r11 factorial r10)
    (mul r13 r11 r1)
    (st tmp r13)
    (ld r14 tmp)
    (ret r14) ) )
(main (n)
      (0 (ld r1 n)
          (call r2 factorial r1)
          (st tmp r2)
          (ld r3 tmp)
          (ret r3) ) ) )
```

Control-Flow Graph Constructions (2 marks)

The design of the internal data structures can significantly impact the extensibility and maintainability of a compiler. Your task is to design a suitable data structure for your optimizer, based on the usual definition of a control-flow graph:

- Nodes in the graph are basic blocks, each consisting of a sequence of instructions. It's fine to have just one or two instructions per block.
- Edges in the graph represent the potential flow of execution.

Take care to take the other tasks of this assignment into consideration when designing the data structure. Document, explain, and justify your design decisions in the relevant code parts.

Unreachable Code (2 marks)

The first optimization eliminates unreachable code. The optimization should be implemented in two phases: first perform a depth-first traversal of the CFG, building a set of all reachable blocks. Remove the unreachable blocks in a second phase.

Dead Code Elimination (3 marks)

Dead code elimination is responsible of cleaning up useless code – that is often left-over from optimization performed before (see below). The optimization consists of a data-flow analysis and a transformation phase.

The data-flow analysis tracks for each register whether it will be *used* subsequently. This is done by performing a backward data-flow analysis, i.e., information is propagated backwards with respect to the flow of execution. A register is considered to be *used* if its value (1) is returned from the current function (`ret`), (2) is stored to a variable (`st`), (3) is used as a condition for a branch (`br`), or (4)

is used by an instructions whose result is used. Your task is to define a data-flow analysis, including lattice, transfer functions, and merge operator. In addition, you should develop an iterative solver that computes a fixed-point solution. Document each component of your analysis thoroughly.

The transformation phase then removes all instructions from the program whose result is not used according to the data-flow analysis.

Example:

```
(0  (ld r1 x)
    (br r1 1 2) )
(1  (lc r2 1)
    (add r1 r1 r2)
    (br r1 3 3) )
(2  (ld r2 x)
    (lc r3 2)
    (add r1 r1 r3)
    (br r1 3 3) )
(3  (ld r5 x)
    (ret r5) )
```

Given the code above, the analysis determines that the value returned from the program (`r5`) along with the corresponding load instruction must be retained. In basic blocks 1 and 2 the addition and load constant instructions must be retained since the result of the addition is used as a condition for the respective branches. The same applies for the conditional branch and the load instruction of basic block 0. The load instruction in basic block 2, however, is not used anywhere (neither by a return, branch, store, or any other instruction whose result has to be retained). The load thus can be removed from the program as shown below.

```
(0  (ld r1 x)
    (br r1 1 2) )
(1  (lc r2 1)
    (add r1 r1 r2)
    (br r1 3 3) )
(2  (lc r3 2)
    (add r1 r1 r3)
    (br r1 3 3) )
(3  (ld r5 x)
    (ret r5) )
```

Elimination of Redundant Loads (3 marks)

Avoiding redundant computations is a frequent way of optimizing programs. For this exercise you should develop an optimization that eliminates redundant load instructions (`ld`). Your solution should consist of a data-flow analysis and a corresponding transformation.

The data-flow analysis determines which registers hold the value of a variable. For example, after a load instruction (`ld r, v`), the destination register `r` holds the value of variable `v`. This remains

unchanged until either `r` is redefined or the value of `v` is overwritten by a store (`st`). Your task is to define a corresponding data-flow analysis problem, consisting of a lattice, transfer functions, and a merge operator. Provide detailed documentation and explanations for each of the analysis's components. You should extend the iterative solver of Task 3 to solve the respective data-flow equations. Once all instructions referring to the register defined by a load instruction are rewritten, the load becomes dead code and can be removed using dead code elimination (Task 3).

Hint: The data-flow analysis problem required to solve this task has the structure of a GEN/KILL problem, a simple and frequent form of analysis problems.

Example:

```
(0  (ld r1 x)
    (br r1 1 2) )
(1  (lc r2 1)
    (add r1 r1 r2)
    (br r1 3 3) )
(2  (ld r2 x)
    (lc r3 2)
    (add r1 r2 r3)
    (br r1 3 3) )
(3  (ld r5 x)
    (ret r5) )
```

Given the basic blocks from above, the analysis determines that at the end of basic block 0, `r1` holds the value of `x`. Basic block 1 does not contain any load, however, register `r1` is redefined. The analysis determines that `x` is not available in any register at the end of basic block 1. Basic block 2 similarly contains a redefinition of register `r1`. In addition the block contains a load instruction, which allows the analysis to determine that `r2` holds the value of `x` at the end of this block. Note that both, `r1` and `r2`, hold the value of `x` up to the point of the redefinition of `r1` within basic block 2. The analysis then determines that `x` is not available in any register at the beginning of basic block 3 by combining the information of basic block 1 and 2.

Using this information, the registers in the program can be rewritten (as shown by the code below). In this example only the use of register `r2` of the addition instruction within basic block 2 is updated to refer to `r1`. This renders the preceding load to `r2` in the same basic block dead code (which will be removed later on).

```
(0  (ld r1 x)
    (br r1 1 2) )
(1  (lc r2 1)
    (add r1 r1 r2)
    (br r1 3 3) )
(2  (ld r2 x)                                -- becomes dead code
    (lc r3 2)
    (add r1 r1 r3)                            -- use of r2 rewritten to r1
    (br r1 3 3) )
(3  (ld r5 x)
    (ret r5) )
```


Implementation

Your assignment should run on the `ucpu[01]` machines. You need to provide all code required to build and execute your optimizer. You need to provide at a minimum a script, `optimize.sh` and a `README` file. The `README` should explain where the relevant code corresponding to each of the tasks can be found. Your `optimize.sh` script should take at least two arguments: the first argument is the name of a file containing unoptimized intermediate code, while the second argument is the name of the intermediate code file that should be produced by your optimizer. Subsequent arguments can be used to select which specific optimization is run. Your `optimize.sh` script should print the available optimizations when run with no arguments.

Further Reading

If you need more background on data-flow analysis, please consult any of the many text books on compiler construction. In particular, the book “Compilers” by Aho, Lam, Sethi, and Ullman provides a very good introduction to the topic.